Exploring Regular Expression Feature Usage in Practice and the Impact on Tool Design

Carl Chapman and Kathryn T. Stolee
Department of Computer Science
Iowa State University
{carl1978, kstolee}@iastate.edu

Abstract—Regular expressions are used frequently in programming languages for form validation, ad-hoc file searches, and simple parsing. Due to the popularity and pervasive use of regular expressions, researchers have created tools to support their creation, validation, and use. Each tool has made design decisions about which regular expression features to support, yet, there does not exist an empirical study of regular expression feature usage to inform these design decisions.

In this paper, we explore regular expression feature usage, focusing on how often features are used and the diversity of regular expressions from syntactic and semantic perspectives. To do this, we analyzed 3898 open source Python projects from GitHub. Our results indicate that the most commonly used regular expression features are also supported by popular research tools and that programmers frequently reinvent the wheel by writing identical or nearly identical regular expressions in different ways.

I. INTRODUCTION

In essence, regular expressions are search patterns for strings. Regular expressions are used extensively in many programming languages, for example, to search text files [3], in form validation [], and for XYZ []. Due in part to their pervasive use across programming languages, many researchers and practitioners have developed tools to support the creation [], validation [], and testing [] of regular expressions.

In writing tools to support regular expressions, tool designers make decisions about which features to support and which not to support. These decisions are often made casually and are dependent on the regular expressions the designers happen to have experience with, the designers have seen in the wild, or their complexity. The goal of this work is to bring more context and information about regular expression feature usage so these decisions can be better informed.

To motivate the study of regular expressions in general, we scrape GitHub for Python projects that use the `re module`, which is the regular expression library for Python. We measure how frequently regular expressions appear in projects, and after doing a feature analysis (e.g., kleene star, literals, and capture groups are all features), we further measure how often such features appear in regular expressions and in projects. Then, we compare the features to those supported by four common regex support tools, brice [], hampi [], Rex [], and RE2 []. We then explore the features not supported by common tools and explore the impact of omitting those features. Our results indicate that these tools support all of the top six most common features and that some of the omitted features, such as the lazy

quantifier, are used in over 35% of projects containing regular expressions.

The contributions of this work are:

- An empirical analysis of the usage of regular expressions in XYZ open-source Python projects

- An analysis of which features are omitted from common regular expression tools and the impact of ignoring those features

- A discussion on the semantic similarity of regular expressions in practice and identification of opportunities for future work in supporting programmers in writing regular expressions.

The rest of the paper is organized as follows. Section II motivates this work by discussing research in supporting programmers in the use, creation, and validation of regular expressions. Section IV presents the research questions and study setup for exploring regular expressions in the wild. Results are in Section V followed by a discussion in Section VI and conclusion.

II. MOTIVATION AND RELATED WORK

Bugs related to regular expressions are common, resulting in tens of thousands of bug reports [4].

III. MOTIVATION

With regexes, there is a common saying: 'now you have two problems'. A skilled programmer can quickly solve many problems using regular expressions, but they can be hard to understand and maintain[back me up somebody - joke guy]. Regular expression languages enable an irreplaceable search technique used within all kinds of text editors and command line tools. Regexes are also employed in critical missions like mysql injection prevention , malicious packet filtering and web form validation.

Tools like Hampi, Kudzu, brics[1], Microsoft's Rex, Automata and the z3-based QF_VRE projects, all attempt to support modeling some subset of regular expression language features. To improve test coverage for code using regular expressions, and to generate strings from regular expressions for whatever other reasons, projects like Reggae,

---

[1]http://www.brics.dk/automaton/

Rex, JST, regex-tester[2], regldg[3], uttool[4], xeger[5], Generex[6], Hoa/Regex[7], Genex[8], Randexp[9], txt2re[10], Pex[11] have been developed. Going the other direction, many projects attempt to take a set of strings and generate a good regular expression that matches them like RegexGenerator++[12], ,Regex-PreSuf[13]. Countless research papers have focused on various other regular expression-related investigations, too many to list. .

One common misconception is that all regular expression languages can be represented using finite state automata. Some regular expression languages

But most regular expression languages are not just regular languages that can be represented by finite state automata anymore, they have . In many regular expression languages

A. Research on Regular Expressions

Visual debugging of regular expressions [1]

Static analysis to reduce errors in building regular expressions by using a type system to identify errors like `PatternSyntaxExceptions` and `IndexOutOfBoundsExceptions` at compile time [4].

B. Research on Regular Expressions

Visual debugging of regular expressions [1]

C. Research that Depends on Regular Expression Usage

Regular expressions are used as queries in a data mining framework [2]

IV. STUDY

To understand how programmers use regular expressions in Python projects (and the syntactic and semantic diversity among the regular expressions), we scraped X projects from GitHub, as described in Section IV-A. Next, we logged all unique regular expressions

We aim to answer the following research questions:

RQ1: How frequently are regexes used in python projects?

To address this question, we measure how often calls to the `re module` are made per file and per project in Python projects.

RQ2: How is the `re module` **TODO: this font looks a little too big - can we make it smaller?** used in python projects?

To address this research question, we measure the frequency of usage for calls to the 8 functions **TODO: link to the diagram?** of the re module(`re.compile`, `re.search`, `re.match`, `re.split`, `re.findall`, `re.finditer`, `re.sub` and `re.subn`) in Python projects scraped from GitHub.

RQ3: Which regex language features are used most commonly in python?

Regex features are components of the regex language, such as capture groups, literals, and the kleene star. To measure feature usage, we parse Python regular expression patterns using Bart Kiers' PCRE parser, as described in Section IV-A.

RQ4: What is the impact of *not* supporting various regex features on tool designers and users?

TODO: clean this up Use semantic analysis to illustrate the impact of missing features on a tool's applicability. Since our semantic analysis is based on Rex, we use syntactic analysis to observe the impact of not supporting various features on this, and other, research.

We map the regex features to each of four research tools that are commonly used for regular expressions research. To address this research question, we looked at the most popular regular expression features that Rex does not support. As Rex is used for our semantic analysis in RQ4, we were interested in the impact of not supporting these features.

TODO: need to justify why we chose the tools we did

A. Building the Corpus

To build a large corpus of regular expressions strings for analysis, we turn to GitHub, a popular project hosting site containing over 100,000 Python projects. We used the GitHub api to page through all repositories, cloning projects that contain Python code. For each project, we used Astroid[X] to build the AST of each Python file and find uses of Python's `re` module. Here is an example of one regex *usage*, with key components labeled:

Fig. 1. example of one regex usage

In the end, we scraped XYZ projects that contained at least some Python code. Of these, XYZ projects also contained at least one Python file with a regular expression, forming the final set of projects for our analysis.

Within each relevant project, duplicate usages (same function, pattern and flags) within the same file (same relative path) were ignored. **TODO: why were duplicates ignored?** Using git, each project was scanned at 20 evenly-spaced commits (or all commits if there were less than 20) in its history. We observed and recorded 53,894 regex usages in 3,898 projects.

Throughout the rest of the paper, we use the following definitions:

Usage: A regex *usage* is a single instance of a regular expression found in a project file.

---

[2] https://github.com/nickawatts/regex-tester

[3] http://regldg.com/

[4] http://uttool.com/text/regexstr/default.aspx

[5] https://code.google.com/p/xeger/

[6] https://github.com/mifmif/Generex

[7] https://github.com/hoaproject/Regex/

[8] http://search.cpan.org/~bowmanbs/Regexp-Genex-0.07/lib/Regexp/Genex.pm

[9] https://www.ruby-toolbox.com/projects/randexp

[10] http://txt2re.com/

[11] http://research.microsoft.com/en-us/projects/pex/

[12] http://regex.inginf.units.it/

[13] http://search.cpan.org/~jhi/Regex-PreSuf-1.17/PreSuf.pm

Pattern: A regex *pattern* is a unique regular expression found in the study and may represent many usages.

That is, the set of patterns is collected by collapsing on duplicate usages. We use patterns as opposed to usages because **TODO: Carl: why is this?**

B. Selecting Patterns

Our analysis focuses on the patterns found, so we ignored the 12.7% of usages using flags that can alter regex behavior. An additional 6.5% of usages contained patterns that could not be compiled because the pattern was non-static (used some runtime variable), or because of other unknown parsing failures.

The remaining 80.8% (43,525) usages were collapsed into 14,113 distinct pattern strings. The resulting set of patten strings were parsed using an antlr-based, open source PCRE parser released by Bart Kiers[14]. This parser was unable to support 0.5% (76) of the patterns due to unsupported unicode characters. Another 0.2% (27) of the patterns used regex features that we have chosen to exclude in this study[15]. The 13,912 distinct pattern strings that remain were each assigned a weight value equal to the number of distinct projects the pattern appeared in. We will refer to this set of weighted, distinct pattern strings as the *collection*.

C. Analyzing Features

After picking four large regex research projects, the big table with the features was created in order to decide which unsupported features are used most often. Our semantic analysis is dependent on the use of Rex to generate strings so we can identify semantically related clusters. For three common features unsupported by Rex, we rely on syntactic analysis to determine similarity among regular expressions containing those features. For those features supported by Rex, we cluster the regular expressions based on semantic diversity.

1) Syntactic Diversity: For the negative perspective, we picked three features: LZY, NCG, WNW that are unsupported by Rex and other projects. For each of these features, we created a subset of the *collection* where all the patterns contain that feature. Then we used syntactic analysis...to create a similarity matrix. We then used markov clustering [X] (MCL) to find clusters in the subset. We used these clusters to assist our manual search for some common use cases for the unsupported feature.

2) Semantic Diversity: For the positive perspective, we created another subset of patterns (XYZ patterns) where Rex was able to generate strings that the pattern matched. We then created a similarity graph with weighted, undirected edges as shown in Figure 2.

Again we used MCL to find clusters that aided a manual search for use cases strongly associated with particular features.

[14]https://github.com/bkiers/pcre-parser
[15]www.details.#thistopic

```
for each row i:
    obtain set of Rex-generated strings Ri from
pattern at index i
    sRi = size of Ri
    for each col j:
        Nij = number of strings in Ri matched by
pattern at index j
        M[i][j] = Nij/sRi
G = empty graph
for each row i:
    for each col j:
        SIMij = (M[i][j]+M[j][i])/2
        if SIMij ¿ 0.75:
            add edge (i,j)=SIMij to G
```
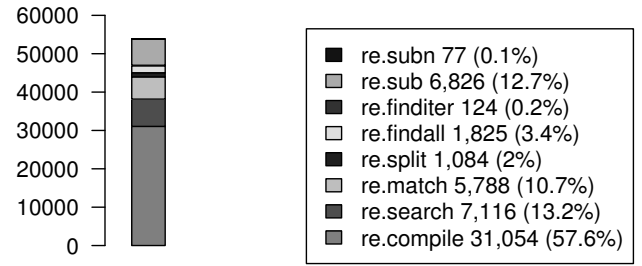
Fig. 2.   Constructing Similarity Graph

Fig. 3.   How often are the 8 re functions used? (RQ2)

V.   RESULTS

A. Context and Corpus

1) Saturation: Although 42.2% of the projects observed had at least one regex usage, only 11.2% of the files observed had at least one regex usage.

From the above figure/table, we see that on average each project had 2 files containing any regex usage, out of an average of 6 files. Each of the files that did have a regex usage had an average of 1 regex usages. Because we scanned 3,898 projects, we would expect to have seen 23,388 regex usages, which is lower than the actual 53,894 usages observed.

2) Regex Functions and Flags: As seen in Figure 3 The 'compile' function encompasses 57.6% of all usages, even though every compiled regex object can only be used by calling other functions. (TODO-Why?)
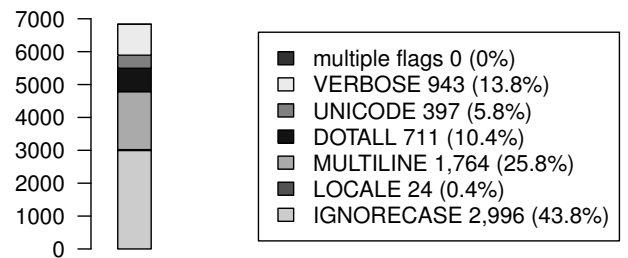

Fig. 4.   Which behavioral flags are used? (RQ2)

87.3% of all regex usages did not use a flag or specified a non-behavioral flag (default or debug). Of all behavioral flags used, ignorecase (43.8%) and multiline (25.8%) were the most frequently used. It is also worth noting that although multiple flags can be combined using a bitwise or, this was never observed. (remove this last part if it is observed later)

3) General Characteristics of Regexes Found: ...TODO

4) Top 10 Regex Patterns by weight:

5) All Features: Literal tokens were found in (TODO) 101% of patterns, and accounted for 75% of all tokens. Excluding literal tokens and features that were not present in any pattern, the following stats...make a sentence, these are some stats about the features:

some more text, IDK

pair	example from corpus	nTimes
CG::ADD	`'(:+)'`	4189
CG::KLE	`'(:)*'`	3983
ANY::KLE	`'.*'`	3709
CG::ANY	`'(.)'`	3160
CCC::CG	`"([')])"`	2665
CCC::ADD	`'[ ]+'`	2612
RNG::CCC	`'[A-Z]'`	2567
ADD::KLE	`'-*(.+)'`	2476
WSP::KLE	`'\\s*'`	2207
END::STR	`'^$'`	2156

OK now that is all for section 2. Now in section 3 I want to look at clustering by string similarity using mcl clustering algorithm. Here are the top 6 clusters using various string similarity metrics:

TODO - multiple boxplots for all 5-6 demonstrating cluster size and then also have # of clusters, pick smallest number of clusters and then use that.

VI. Discussion

...only 11.2% of the files observed had at least one regex usage. This indicates that regex usage may usually be concentrated in just a few files.

Fun fact: while creating similarity matrix, row 5464 took 2 hours, or almost 1 second per cell avg, only suffering 18 timeouts (1.2 secs). What is this pesky pattern?

VII. Conclusion

Acknowledgment

References

[1] F. Beck, S. Gulan, B. Biegel, S. Baltes, and D. Weiskopf. Regviz: Visual debugging of regular expressions. In *Companion Proceedings of the 36th International Conference on Software Engineering*, ICSE Companion 2014, pages 504–507, New York, NY, USA, 2014. ACM.

[2] A. Begel, Y. P. Khoo, and T. Zimmermann. Codebook: Discovering and exploiting relationships in software repositories. In *Proceedings of the 32Nd ACM/IEEE International Conference on Software Engineering - Volume 1*, ICSE '10, pages 125–134, New York, NY, USA, 2010. ACM.

[3] C. L. A. Clarke and G. V. Cormack. On the use of regular expressions for searching text. *ACM Trans. Program. Lang. Syst.*, 19(3):413–426, May 1997.

[4] E. Spishak, W. Dietl, and M. D. Ernst. A type system for regular expressions. In *Proceedings of the 14th Workshop on Formal Techniques for Java-like Programs*, FTfJP '12, pages 20–26, New York, NY, USA, 2012. ACM.

pattern	weight
`'\\s+'`	181
`'\\s'`	78
`'\\d+'`	70
`'[\\x80-\\xff]'`	69
`'\nmd5_data = {\n([^}]+)}'`	69
`'\\\\(.)'`	67
`'([\\\\"]|[^\\ -~])'`	66
`'(-?(?:0|[1-9]\\d*))(\\.\\d+)?([eE][-+]?\\d+)?'`	61
`'[^]]+?\\] +([0-9.]+): (\\w+) <-(\\w+)'`	60
`'.*rlen=([0-9]+)'`	57

rank	code	description	example	brics	hampi	Rex	RE2	nPatterns	% patterns	nFiles	%files	nProjects	% projects
1	ADD	one-or-more repetition	`z+`	●	●	●	●	6,122	44	9,330	50.3	1,209	73.5
2	CG	a capture group	`(caught)`	●	●	●	●	7,248	52.1	9,759	52.6	1,197	72.8
3	KLE	zero-or-more repetition	`.*`	●	●	●	●	6,104	43.9	8,323	44.9	1,100	66.9
4	CCC	custom character class	`[aeiou]`	●	●	●	●	4,581	32.9	7,808	42.1	1,027	62.4
5	ANY	any non-newline char	`.`	●	●	●	●	4,708	33.8	6,394	34.5	1,006	61.2
6	RNG	chars within a range	`[a-z]`	●	●	●	●	2,698	19.4	5,196	28	849	51.6
7	STR	start-of-line	`^`	○	●	●	●	3,660	26.3	5,622	30.3	847	51.5
8	END	end-of-line	`$`	○	●	●	●	3,258	23.4	5,549	29.9	828	50.3
9	NCCC	negated CCC	`[^qwxf]`	●	●	●	●	1,970	14.2	4,027	21.7	777	47.2
10	WSP	\t \n \r \b \f or space	`\s`	○	●	●	●	2,908	20.9	4,812	25.9	764	46.4
11	OR	logical or	`a|b`	●	●	●	●	2,161	15.5	4,039	21.8	711	43.2
12	DEC	any of: 0123456789	`\d`	○	●	●	●	2,385	17.1	4,366	23.5	694	42.2
13	WRD	[a-zA-Z0-9_]	`\w`	○	●	●	●	1,457	10.5	3,004	16.2	652	39.6
14	QST	zero-or-one repetition	`z?`	●	●	●	●	1,922	13.8	3,821	20.6	647	39.3
15	LZY	as few reps as possible	`z+?`	○	●	○	●	1,318	9.5	2,291	12.4	606	36.8
16	NCG	group without capturing	`a(?:b)c`	○	●	○	●	813	5.8	1,748	9.4	404	24.6
17	NCG	named capture group	`(?P<name>x)`	○	●	○	●	934	6.7	1,517	8.2	354	21.5
18	SNG	exactly n repetition	`z{8}`	●	●	●	●	623	4.5	1,359	7.3	340	20.7
19	NWSP	any non-whitespace	`\S`	○	●	●	●	490	3.5	788	4.2	271	16.5
20	DBB	$n \le x \le m$ repetition	`z{3,8}`	●	●	●	●	384	2.8	692	3.7	242	14.7
21	NLKA	sequence doesn't follow	`a(?!yz)`	○	○	●	○	137	1	503	2.7	184	11.2
22	NWRD	non-word chars	`\W`	○	●	●	●	97	0.7	315	1.7	169	10.3
23	LWB	at least n repetition	`z{15,}`	●	●	●	●	97	0.7	337	1.8	167	10.2
24	WNW	word/non-word boundary	`\b`	○	○	○	●	248	1.8	438	2.4	166	10.1
25	LKA	matching sequence follows	`a(?=bc)`	○	○	○	○	114	0.8	360	1.9	159	9.7
26	OPT	options wrapper	`(?i)CasE`	○	●	○	●	232	1.7	378	2	154	9.4
27	NLKB	sequence doesn't precede	`(?<!x)yz`	○	○	○	○	102	0.7	321	1.7	139	8.4
28	LKB	matching sequence precedes	`(?<=a)bc`	○	○	○	○	82	0.6	262	1.4	120	7.3
29	ENDZ	absolute end of string	`\Z`	○	○	○	●	91	0.7	154	0.8	94	5.7
30	BKR	match the i^{th} CG	`\1`	○	○	○	○	60	0.4	129	0.7	84	5.1
31	NDEC	any non-decimal	`\D`	○	●	●	●	36	0.3	92	0.5	58	3.5
32	BKRN	references NCG	`\g<name>`	○	●	○	○	17	0.1	44	0.2	28	1.7
33	VWSP	matches U+000B	`\v`	○	○	●	●	13	0.1	16	0.1	15	0.9
34	NWNW	negated WNW	`\B`	○	○	○	●	4	0	11	0.1	11	0.7